ML Week

Natural Language Processing

Jeff Abrahamson

23–24 novembre 2016

Linear Programming

$$\begin{aligned} &\text{Maximize} && c^T x \\ &\text{subject to} && Ax \leqslant b \end{aligned}$$

Summarising Text

- Abstractive (hard)
- Extractive (select sentences)

Summarising Text

Challenge problem (cf. greedy solutions):

The cat is in the kitchen.

The cat drinks the milk.

The cat drinks the milk in the kitchen.

Summarising Text

- Sentence selection
- Use n-grams
- Stemming
- Stop words
- Prune short sentences

Dan Gillick, Benoit Favre, A Scalable Global Model for Summarization, 2009

Summarising Text

Outline:

- ILP *(optimisation linéaire en nombres entiers)*
- Maximum coverage model

Dan Gillick, Benoit Favre, A Scalable Global Model for Summarization, 2009

Summarising Text

ILP in canonical form:

$$\begin{aligned} \text{Maximize} \quad & c^T x \\ \text{subject to} \quad & Ax \leqslant b \\ & x \geqslant 0 \\ & x \in \mathbb{Z}^n \end{aligned}$$

Dan Gillick, Benoit Favre, A Scalable Global Model for Summarization, 2009

Summarising Text

ILP in standard form:

$$\begin{aligned} \text{Maximize} \quad & c^T x \\ \text{subject to} \quad & Ax + s = b \\ & s \geqslant 0 \\ & x \in \mathbb{Z}^n \end{aligned}$$

Dan Gillick, Benoit Favre, A Scalable Global Model for Summarization, 2009

Summarising Text

ILP in standard form:

$$\begin{aligned} \text{Maximize} \quad & c^T x \\ \text{subject to} \quad & Ax + s = b \\ & s \geqslant 0 \\ & x \in \mathbb{Z}^n \end{aligned}$$

This is NP hard.

Dan Gillick, Benoit Favre, A Scalable Global Model for Summarization, 2009

Summarising Text

ILP in standard form:

$$\begin{aligned} \text{Maximize} \quad & c^T x \\ \text{subject to} \quad & Ax + s = b \\ & s \geqslant 0 \\ & x \in \mathbb{Z}^n \end{aligned}$$

Discussion: linear vs integer programming.

Dan Gillick, Benoit Favre, A Scalable Global Model for Summarization, 2009

Summarising Text

Let

$$
\begin{aligned}
c_i &: \text{presence of concept } i \text{ in summary} \\
w_i &: \text{weight associated with } c_i \\
l_i &: \text{length of sentence } i \\
s_j &: \text{presence of sentence } j \text{ in summary} \\
L &: \text{summary length limit} \\
Occ_{ij} &: \text{occurence of } c_i \text{ in } s_j
\end{aligned}
$$

Dan Gillick, Benoit Favre, A Scalable Global Model for Summarization, 2009

Summarisation

$$
\begin{aligned}
\text{Maximize} \quad & \sum_i w_i c_i \\
\text{subject to} \quad & \sum_j l_j s_j \leqslant L \\
& s_j Occ_{ij} \leqslant c_i, \qquad \forall i, j \\
& \sum_j s_j Occ_{ij} \geqslant c_i \qquad \forall i \\
& c_j \in \{0, 1\}, \qquad \forall i \\
& s_j \in \{0, 1\}, \qquad \forall j
\end{aligned}
$$

Dan Gillick, Benoit Favre, A Scalable Global Model for Summarization, 2009

Summarising Text

Notes:

- Selecting a sentence selects all concepts it contains
- Selecting a concept requires it be in at least one sentence
- $s_j Occ_{ij} \leqslant c_i, \ \forall i, j \ \Rightarrow$ no concept-less sentences

Dan Gillick, Benoit Favre, A Scalable Global Model for Summarization, 2009

Sentiment Analysis

Many variations:

- Entire documents using computational linguistics
- Manually crafted lexicons

Sentiment Analysis

Techniques

- Template instantiation (requires domain knowledge)
- Passage extraction

Sentiment Analysis

- Extract “opinion sentences” based on the presence of a predetermined list of product features and adjectives.
- Evaluate the sentences based on counts of positive vs negative polarity words (as determined by the Wordnet algorithm)

Hu and Lieu, Mining and Summarizing Customer Reviews, 2004

Sentiment Analysis

- Extract “opinion sentences” based on the presence of a predetermined list of product features and adjectives.
  - “The food is excellent.”
  - “The food is an excellent example of how not to cook.”
- Evaluate the sentences based on counts of positive vs negative polarity words (as determined by the Wordnet algorithm)

Hu and Lieu, Mining and Summarizing Customer Reviews, 2004

Sentiment Analysis

- Extract “opinion sentences” based on the presence of a predetermined list of product features and adjectives.
- Evaluate the sentences based on counts of positive vs negative polarity words (as determined by the Wordnet algorithm)

The good: fast, no training data, decent prediction.

The bad: fails on multiple word sense, non-adjectives; sensitive to context.

Hu and Lieu, Mining and Summarizing Customer Reviews, 2004

Sentiment Analysis

Words aren't enough.

- "unpredictable plot" vs "unpredictable performance"

Turney, Thumbs Up or Thumbs Down? Semantic Orientation Applied to Unsupervised Classification of Reviews, 2002

Questions?

ml-week.com/1